

Databases for Analytics

Basic SQLite Usage

Learning Objectives

- **Skills:** You should know how to ...
 - Use the `sqlite3` command line interpreter (CLI)
 - Create a SQLite database from scratch from the CLI and from Python (Jupyter)
 - Write SQLite-flavored SQL scripts
 - Populate the database from MySQL or CSV files
- **Theory:** You should be able to explain ...
 - SQLite's in-memory database mode
 - The pros and cons of SQLite for analytics

SQLite FTW!

SQLite does not require a DBMS or an installer. Any computer with a modern version of Python can use SQLite to store relational data.

- **File-based**, so that we can copy files and use them directly (and even check them into GitHub)
- **Tiny**, requiring almost no CPU time or memory
- **(Mostly) Standard SQL**, so we can use what we already know

SQLite from the Command Line

The old-school way to use SQLite

The `sqlite3` Interpreter

While SQLite is not a full-featured DBMS, the `sqlite3` interpreter does many of the same functions.

- Create a new database / Open an existing one
- Issue SQL DDL and DML
- Inspect database schema (tables, indexes, etc.)
- Dump/load to/from SQL files

CLI Docs: <https://www.sqlite.org/cli.html>

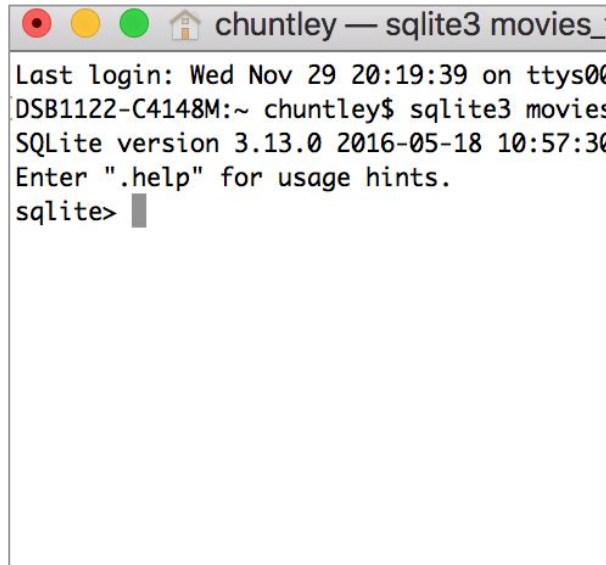
Starting `sqlite3`

`sqlite3` is included in JupyterLab
(and in MacOS)

Step 1. Git → Open Terminal to open a terminal window in your repo folder.

Step 2. Type **`sqlite3 database.db`**

- **`database.db`** is just a file path
- Creates the file if needed
- If no filename, then it will use in memory mode

A terminal window titled "chuntley — sqlite3 movies_" with standard macOS window controls (red, yellow, green buttons). The terminal text shows the command "sqlite3 movies_" being executed, followed by the SQLite version information and a prompt "sqlite>".

```
chuntley — sqlite3 movies_
Last login: Wed Nov 29 20:19:39 on ttys00
DSB1122-C4148M:~ chuntley$ sqlite3 movies_
SQLite version 3.13.0 2016-05-18 10:57:30
Enter ".help" for usage hints.
sqlite>
```

In-Memory vs File Storage

SQLite is designed to be very lightweight, able to work on even the smallest devices.

In fact, it does not need need a proper file system! We can use SQLite in memory, with any permanent storage at all.

Just fire up `sqlite3` without specifying a filename.

Dot Commands

Admin functions are available using dot commands. To get a list just use the **.help** command.

Note that the **.exit** command is used to close `sqlite3`.

```
chuntley — sqlite3 movies_tonight.db — 82x35

[sqlite> .help
.auth ON|OFF          Show authorizer callbacks
.backup ?DB? FILE     Backup DB (default "main") to FILE
.bail on|off          Stop after hitting an error. Default OFF
.binary on|off        Turn binary output on or off. Default OFF
.changes on|off       Show number of rows changed by SQL
.clone NEWDB          Clone data into NEWDB from the existing database
.databases            List names and files of attached databases
.dbinfo ?DB?         Show status information about the database
.dump ?TABLE? ...    Dump the database in an SQL text format
                     If TABLE specified, only dump tables matching
                     LIKE pattern TABLE.

.echo on|off          Turn command echo on or off
.eqp on|off|full      Enable or disable automatic EXPLAIN QUERY PLAN
.exit                Exit this program
.explain ?on|off|auto? Turn EXPLAIN output mode on or off or to automatic
.fullschema ?--indent? Show schema and the content of sqlite_stat tables
.headers on|off       Turn display of headers on or off
.help                Show this message
.import FILE TABLE   Import data from FILE into TABLE
.indexes ?TABLE?      Show names of all indexes
                     If TABLE specified, only show indexes for tables
                     matching LIKE pattern TABLE.

.limit ?LIMIT? ?VAL? Display or change the value of an SQLITE_LIMIT
.load FILE ?ENTRY?    Load an extension library
.log FILE|off         Turn logging on or off. FILE can be stderr/stdout
.mode MODE ?TABLE?    Set output mode where MODE is one of:
                     ascii  Columns/rows delimited by 0x1F and 0x1E
                     csv    Comma-separated values
                     column Left-aligned columns. (See .width)
                     html   HTML <table> code
                     insert  SQL insert statements for TABLE
                     line   One value per line
                     list    Values delimited by .separator strings
                     tabs    Tab-separated values
```


Inspecting Database Schema

.tables

- Lists all the tables in the database
- Equivalent to the MySQL show tables command

.indexes

.indexes *tablename*

- Can be used to see every PK, FK, or other index in a table

.schema

.schema *tablename*

- Shows the CREATE TABLE DDL for one or more tables

CSV Import/Export

Importing from a CSV file:

```
sqlite> .mode csv  
sqlite> .import filename.csv tablename
```

If the table already exists then the columns headers are treated as data. It's usually best to make sure the table is empty.

Exporting to a CSV file:

```
sqlite> .header on  
sqlite> .mode csv  
sqlite> .once filename.csv  
sqlite> SELECT * FROM tablename;
```

You'll likely want to set **.mode list** afterwards.

.once is used to direct query output to the file one time and then switch output back to the screen.

Dumping / Loading the Database

.dump

- Need to call **.once** first to direct the dump to a file.
- The SQL DDL is just a direct copy of whatever was typed in, capitalization and spacing included.

.read *filename*

- Loads and executes the SQL in the file against the current database.
- Works best with file produced by the **.dump** command

SQLite in Jupyter

A more lightweight and less CLI-heavy approach

TL;DR: Use `%sql` where you can

As long as we abide by the minor SQL dialect differences, SQLite queries in Jupyter work pretty much the same as they did with MySQL, only simpler.

```
%load_ext sql
%sql sqlite:///deals.db
c = %sql SELECT * FROM COMPANIES
companies = c.DataFrame()
```

PRAGMA

If you need to inspect tables or other SQLite metadata programmatically (e.g., in a Jupyter Notebook), use the special PRAGMA pseudo-SQL statement.

```
PRAGMA table_info(tablename);
```

```
PRAGMA foreign_key_list(tablename);
```

Docs: <https://www.sqlite.org/pragma.html>

Connection Strings

Since SQLite does not have username or passwords and the database file is always on the local computer, the usual SQLAlchemy connection string reduces to `sqlite:///filename`

Note that that's **3 slashes** before the *filename*.

Databases for Analytics

Basic SQLite Usage